



## Department of Computer Science and Engineering (Data Science)

**Subject: Image Processing and Computer Vision - II Laboratory (DJ19DSL702)**

**AY: 2024-25**

### **Experiment 7**

#### **( Object Segmentation Using U-Net)**

Name: Riya Chavan

Sap: 60009220136

Batch: D2-1

**Aim:** To segment objects in an image using the U-Net architecture.

#### **Theory:**

##### **1. Introduction**

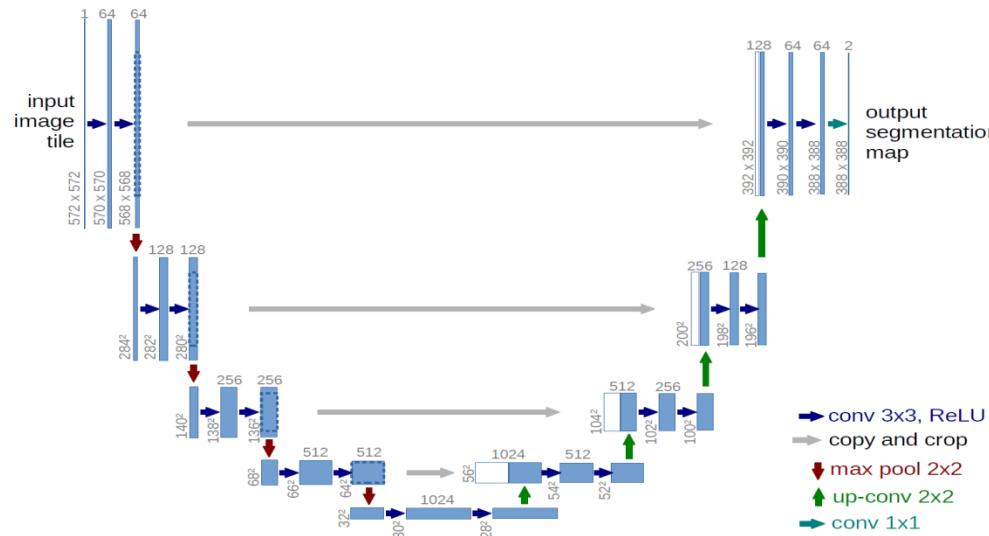
U-Net is known for its encoder-decoder structure and skip connections that help in segmenting biomedical and general images with limited data. The u-net is convolutional network architecture for fast and precise segmentation of images. Up to now it has outperformed the prior best method (a sliding-window convolutional network) on the ISBI challenge for segmentation of neuronal structures in electron microscopic stacks. It has won the Grand Challenge for Computer-Automated Detection of Caries in Bitewing Radiography at ISBI 2015, and it has won the Cell Tracking Challenge at ISBI 2015 on the two most challenging transmitted light microscopy categories (Phase contrast and DIC microscopy) by a large margin

##### **2. Architecture**

The U-Net includes a contracting path (encoder), expansive path (decoder), and skip connections that help retain spatial details.



## Department of Computer Science and Engineering (Data Science)



### Encoder Network

The encoder network acts as the feature extractor and learns an abstract representation of the input image through a sequence of the encoder blocks. Each encoder block consists of two 3x3 convolutions, where each convolution is followed by a ReLU (Rectified Linear Unit) activation function. The ReLU activation function introduces non-linearity into the network, which helps in the better generalization of the training data. The output of the ReLU acts as a skip connection for the corresponding decoder block. Next, follows a 2x2 max-pooling, where the spatial dimensions (height and width) of the feature maps are reduced by half. This reduces the computational cost by decreasing the number of trainable parameters.

### Skip Connections

These skip connections provide additional information that helps the decoder to generate better semantic features. They also act as a shortcut connection that helps the indirect flow of gradients to the earlier layers without any degradation. In simple terms, we can say that skip connection helps in better flow of gradient while backpropagation, which in turn helps the network to learn better representation.



## Bridge

The bridge connects the encoder and the decoder network and completes the flow of information. It consists of two 3x3 convolutions, where each convolution is followed by a ReLU activation function.

## Decoder Network

The decoder network is used to take the abstract representation and generate a semantic segmentation mask. The decoder block starts with a 2x2 transpose convolution. Next, it is concatenated with the corresponding skip connection feature map from the encoder block. These skip connections provide features from earlier layers that are sometimes lost due to the depth of the network. After that, two 3x3 convolutions are used, where each convolution is followed by a ReLU activation function.

The output of the last decoder passes through a 1x1 convolution with sigmoid activation. The sigmoid activation function gives the segmentation mask representing the pixel-wise classification.

## Dataset Preparation

Use the same dataset as PSPNet with masks one-hot encoded for multi-class segmentation.

## Implementation Steps

1. Preprocess the dataset: normalize images and prepare segmentation masks.
2. Build the U-Net model with encoder-decoder blocks and skip connections.
3. Choose Binary Cross-Entropy or Categorical Cross-Entropy as the loss function.
4. Use Adam optimizer with learning rate scheduling.



## Department of Computer Science and Engineering (Data Science)

5. Train and validate the model using accuracy and loss metrics.
6. Test the model and compare predictions with ground truth masks.

### Result and Evaluation

Use IoU, and visualization plots to assess model performance.

### Conclusion

U-Net effectively segments images by combining localization and context through skip connections.

```
# Import necessary libraries
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.models as models
from torchvision import transforms
from torch.utils.data import Dataset, DataLoader
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# Define U-Net architecture components
class DoubleConv(nn.Module):
    """(Conv => BN => ReLU) * 2"""
    def __init__(self, in_channels, out_channels, mid_channels=None):
        super().__init__()
        if not mid_channels:
            mid_channels = out_channels
        self.double_conv = nn.Sequential(
            nn.Conv2d(in_channels, mid_channels, kernel_size=3, padding=1,
bias=False),
            nn.BatchNorm2d(mid_channels),
            nn.ReLU(inplace=True),
            nn.Conv2d(mid_channels, out_channels, kernel_size=3,
padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True)
        )
```



## Department of Computer Science and Engineering (Data Science)

```
def forward(self, x):
    return self.double_conv(x)

class Down(nn.Module):
    """Downscaling with maxpool then double conv"""
    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.maxpool_conv = nn.Sequential(
            nn.MaxPool2d(2),
            DoubleConv(in_channels, out_channels)
        )

    def forward(self, x):
        return self.maxpool_conv(x)

class Up(nn.Module):
    """Upscaling then double conv"""
    def __init__(self, in_channels, out_channels, bilinear=True):
        super().__init__()

        # if bilinear, use the normal convolutions to reduce the number of
        channels
        if bilinear:
            self.up = nn.Upsample(scale_factor=2, mode='bilinear',
align_corners=True)
            self.conv = DoubleConv(in_channels, out_channels, in_channels
// 2)
        else:
            self.up = nn.ConvTranspose2d(in_channels, in_channels // 2,
kernel_size=2, stride=2)
            self.conv = DoubleConv(in_channels, out_channels)

    def forward(self, x1, x2):
        x1 = self.up(x1)
        # input is CHW
        diffY = x2.size()[2] - x1.size()[2]
        diffX = x2.size()[3] - x1.size()[3]

        x1 = F.pad(x1, [diffX // 2, diffX - diffX // 2,
                      diffY // 2, diffY - diffY // 2])
        x = torch.cat([x2, x1], dim=1)
```



## Department of Computer Science and Engineering (Data Science)

```
return self.conv(x)

class OutConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(OutConv, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, kernel_size=1)

    def forward(self, x):
        return self.conv(x)

# Define U-Net model
class UNet(nn.Module):
    def __init__(self, n_channels=3, n_classes=150, bilinear=True):
        super(UNet, self).__init__()
        self.n_channels = n_channels
        self.n_classes = n_classes
        self.bilinear = bilinear

        # Encoder path
        self.inc = DoubleConv(n_channels, 64)
        self.down1 = Down(64, 128)
        self.down2 = Down(128, 256)
        self.down3 = Down(256, 512)
        factor = 2 if bilinear else 1
        self.down4 = Down(512, 1024 // factor)

        # Decoder path with skip connections
        self.up1 = Up(1024, 512 // factor, bilinear)
        self.up2 = Up(512, 256 // factor, bilinear)
        self.up3 = Up(256, 128 // factor, bilinear)
        self.up4 = Up(128, 64, bilinear)
        self.outc = OutConv(64, n_classes)

    def forward(self, x):
        # Encoder path
        x1 = self.inc(x)
        x2 = self.down1(x1)
        x3 = self.down2(x2)
        x4 = self.down3(x3)
        x5 = self.down4(x4)
```



## Department of Computer Science and Engineering (Data Science)

```
# Decoder path with skip connections
x = self.up1(x5, x4)
x = self.up2(x, x3)
x = self.up3(x, x2)
x = self.up4(x, x1)
logits = self.outc(x)

return logits

# Define Dataset Loader with one-hot encoding for masks
class ADE20KDataset(Dataset):
    def __init__(self, root_dir, image_set='training', transform=None,
target_transform=None):
        self.image_dir = os.path.join(root_dir, 'images', image_set)
        self.mask_dir = os.path.join(root_dir, 'annotations', image_set)
        self.images = sorted(os.listdir(self.image_dir))
        self.masks = sorted(os.listdir(self.mask_dir))
        self.transform = transform
        self.target_transform = target_transform
        self.n_classes = 150

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        image = Image.open(os.path.join(self.image_dir,
self.images[idx])).convert('RGB')
        mask = Image.open(os.path.join(self.mask_dir, self.masks[idx]))

        if self.transform:
            image = self.transform(image)

        if self.target_transform:
            mask = self.target_transform(mask)
        else:
            # Default resize to match model output
            mask = mask.resize((256, 256), resample=Image.NEAREST)
            mask = np.array(mask)

        # Set ignore label (150) to 255 for CrossEntropyLoss to ignore
        mask[mask >= 150] = 255
```



## Department of Computer Science and Engineering (Data Science)

```
mask = torch.tensor(mask, dtype=torch.long)

return image, mask

# Define metrics for evaluation
def compute_iou(pred, target, n_classes=150, ignore_index=255):
    """Compute IoU for semantic segmentation"""
    pred = torch.argmax(pred, dim=1)

    # Create mask for valid pixels
    mask = (target != ignore_index)

    ious = []
    for cls in range(n_classes):
        pred_inds = (pred == cls) & mask
        target_inds = (target == cls) & mask

        intersection = (pred_inds & target_inds).sum().item()
        union = (pred_inds | target_inds).sum().item()

        if union == 0:
            ious.append(float('nan')) # Class not present in ground truth
or prediction
        else:
            ious.append(intersection / union)

    return np.nanmean(ious) # Average over classes

# Train the Model
def train_model():
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

    # Data preprocessing
    transform = transforms.Compose([
        transforms.Resize((256, 256)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])
    ])
```



## Department of Computer Science and Engineering (Data Science)

```
# Create datasets
train_dataset = ADE20KDataset(root_dir='ADEChallengeData2016',
image_set='training', transform=transform)
val_dataset = ADE20KDataset(root_dir='ADEChallengeData2016',
image_set='validation', transform=transform)

# Create dataloaders
train_loader = DataLoader(train_dataset, batch_size=8, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=8, shuffle=False)

# Create model
model = UNet(n_channels=3, n_classes=150, bilinear=True).to(device)

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss(ignore_index=255)
optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)

# Learning rate scheduler
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, mode='min', factor=0.5, patience=3, verbose=True
)

# Training history
history = {
    'train_loss': [],
    'train_iou': [],
    'val_loss': [],
    'val_iou': []
}

# Training loop
num_epochs = 3
for epoch in range(num_epochs):
    print(f"\nEpoch {epoch+1}/{num_epochs}")

    # Train
    model.train()
    train_loss = 0
    train_iou = 0
    batch_count = 0
```



### Department of Computer Science and Engineering (Data Science)

```
for images, masks in train_loader:
    images, masks = images.to(device), masks.to(device)

    # Forward pass
    outputs = model(images)
    loss = criterion(outputs, masks)

    # Backward pass
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # Calculate IoU
    iou = compute_iou(outputs, masks)

    train_loss += loss.item()
    train_iou += iou
    batch_count += 1

    if batch_count % 10 == 0:
        print(f"Batch {batch_count}, Loss: {loss.item():.4f}, IoU: {iou:.4f}")

    avg_train_loss = train_loss / batch_count
    avg_train_iou = train_iou / batch_count

    # Validation
    model.eval()
    val_loss = 0
    val_iou = 0
    batch_count = 0

    with torch.no_grad():
        for images, masks in val_loader:
            images, masks = images.to(device), masks.to(device)

            outputs = model(images)
            loss = criterion(outputs, masks)
            iou = compute_iou(outputs, masks)

            val_loss += loss.item()
```



## Department of Computer Science and Engineering (Data Science)

```
        val_iou += iou
        batch_count += 1

    avg_val_loss = val_loss / batch_count
    avg_val_iou = val_iou / batch_count

    # Update learning rate
    scheduler.step(avg_val_loss)

    # Save history
    history['train_loss'].append(avg_train_loss)
    history['train_iou'].append(avg_train_iou)
    history['val_loss'].append(avg_val_loss)
    history['val_iou'].append(avg_val_iou)

    print(f"Epoch {epoch+1}, Train Loss: {avg_train_loss:.4f}, Train
IoU: {avg_train_iou:.4f}")
    print(f"Epoch {epoch+1}, Val Loss: {avg_val_loss:.4f}, Val IoU:
{avg_val_iou:.4f}")

    # Save model
    torch.save(model.state_dict(), "unet_ade20k.pth")
    print("☑ Model trained and saved as unet_ade20k.pth")

    # Plot training history
    plot_training_history(history)

    return model, history

# Visualization functions
def plot_training_history(history):
    """Plot training and validation loss and IoU"""
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))

    # Plot loss
    ax1.plot(history['train_loss'], label='Train')
    ax1.plot(history['val_loss'], label='Validation')
    ax1.set_title('Loss')
    ax1.set_xlabel('Epoch')
    ax1.set_ylabel('Loss')
    ax1.legend()
```



**Department of Computer Science and Engineering (Data Science)**

```
# Plot IoU
ax2.plot(history['train_iou'], label='Train')
ax2.plot(history['val_iou'], label='Validation')
ax2.set_title('IoU')
ax2.set_xlabel('Epoch')
ax2.set_ylabel('IoU')
ax2.legend()

plt.tight_layout()
plt.show()

def visualize_predictions(model, image_path):
    """Visualize model predictions on a single image"""
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    model.eval()

    # Load and preprocess the image
    transform = transforms.Compose([
        transforms.Resize((256, 256)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])
    ])

    image = Image.open(image_path).convert('RGB')
    input_tensor = transform(image).unsqueeze(0).to(device)

    # Generate prediction
    with torch.no_grad():
        output = model(input_tensor)
        pred = torch.argmax(output, dim=1)[0].cpu().numpy()

    # Denormalize image for visualization
    input_np = input_tensor[0].cpu().numpy()
    input_np = np.transpose(input_np, (1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    input_np = input_np * std + mean
    input_np = np.clip(input_np, 0, 1)
```



### Department of Computer Science and Engineering (Data Science)

```
# Create a colormap for visualization
cmap = plt.cm.get_cmap('tab20', 151)

# Plot results
plt.figure(figsize=(10, 5))

plt.subplot(1, 2, 1)
plt.imshow(input_np)
plt.title('Input Image')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(pred, cmap=cmap, vmin=0, vmax=150)
plt.title('Segmentation Prediction')
plt.axis('off')

plt.tight_layout()
plt.show()

model, history = train_model()

# Test on a sample image
visualize_predictions(model,
"/content/ADEChallengeData2016/images/validation/ADE_val_00000015.jpg")
```



## Department of Computer Science and Engineering (Data Science)

